
Synthetic Data Compilation for Rollout-Free Trajectory Planning

Thomas Y.L. Lin¹

Abstract

We revisit model-based reinforcement learning and remove recurrent rollout for downstream planning. Current action-conditioned models confront an efficiency-accuracy tradeoff for prediction: learning immediate next-state transitions requires fine temporal discretization and numerous rollout steps, while accurate regression needs massive trajectory data. We solve this tension by distilling a teacher representation of instantaneous dynamics into a student predictor. This gives both planning and training advantages. For planning, the student maps an action sequence to its consequences, avoiding recurrent rollout and backpropagation through a recurrent graph. For training, the teacher compiles synthetic supervision from any well-posed integration of the learned dynamics, improving generalization across trajectory distributions while preserving consistency with the governing law. We provide theoretical justification and empirical evidence for these advantages.

1 Introduction

We revisit model-based reinforcement learning (MBRL) and remove recurrent rollout for downstream planning. Our motivation comes from mainstream action-conditioned world models (Hafner et al., 2020; Hansen et al., 2022b) that map a current state and action to the immediate next state, implicitly assuming locally constant actions under zero-order hold (ZOH). Such modeling relies on fine temporal discretization and thus requires numerous recurrent rollout steps for planning. For example, with $\Delta t \approx 0.02\text{s}$, a 10-second query requires 500 rollout steps (Tassa et al., 2018; Hafner et al., 2020; Hansen et al., 2022b). Moreover, action planning through this recurrent graph becomes costly and susceptible to vanishing or exploding gradients (Pascanu et al., 2013). Removing recurrent rollout suggests learning an endpoint predictor that maps an initial state and an action sequence directly to its consequence. However, direct endpoint regression is high-dimensional, so trajectory data coverage be-

comes sparse as the action-sequence length grows. Together, rollout-based modeling and endpoint prediction expose a tradeoff between planning cost and large-scale supervision. We resolve this tradeoff by compiling synthetic data from a learned teacher dynamics model. The teacher predicts the instantaneous state derivative from local state-action pairs, i.e. $\dot{z} = f(z, u)$. For each supported action sequence, we integrate the teacher along the induced trajectory and use the resulting endpoint to distill a student predictor.

Advantages. The proposed method (Figure 1) shows three advantages:

- **Remove Planning Rollout.** The endpoint predictor replaces recurrent rollout with direct consequence query. When the reward objective uses K checkpoints (9), the model evaluates K endpoint queries that can be batched in parallel rather than sequentially (Section 3).
- **Direct Action Gradients.** Planning through the predictor optimizes the action sequence by backpropagating through a single network instead of a recurrent one. This removes products of Jacobians and gradient sensitivity in rollout planning (Lemma 3.1 and Corollary 3.2).
- **Synthetic Endpoint Supervision.** We compile synthetic endpoint targets by integrating the teacher along supported induced trajectories. This improves coverage across high-dimensional trajectory distributions while preserving consistency with the governing law (Section 4 and Proposition 4.1).

Roadmap. Section 2 formalizes. Sections 3 and 4 justify the advantages. Sections B and 5 report empirical tests. Section 6 concludes. Section A reviews related work.

2 Endpoint World Model

Controlled dynamical law. For latent state $z(t) \in \mathbb{R}^d$ and control $u(t) \in \mathbb{R}^m$, assume the governing law f follows

$$\frac{dz}{dt} = f(z(t), u(t)).$$

Note this law is autonomous: all relevant exogenous variables are included in z or u , so f has no explicit time dependency. For any well-posed control signal $u|_{[t, t+T]} := u(\cdot)|_{[t, t+T]}$, f induces an *endpoint operator* Φ_T^f mapping

¹University of Washington, Seattle, USA. Correspondence to: Thomas Y.L. Lin <thermaus@uw.edu>.

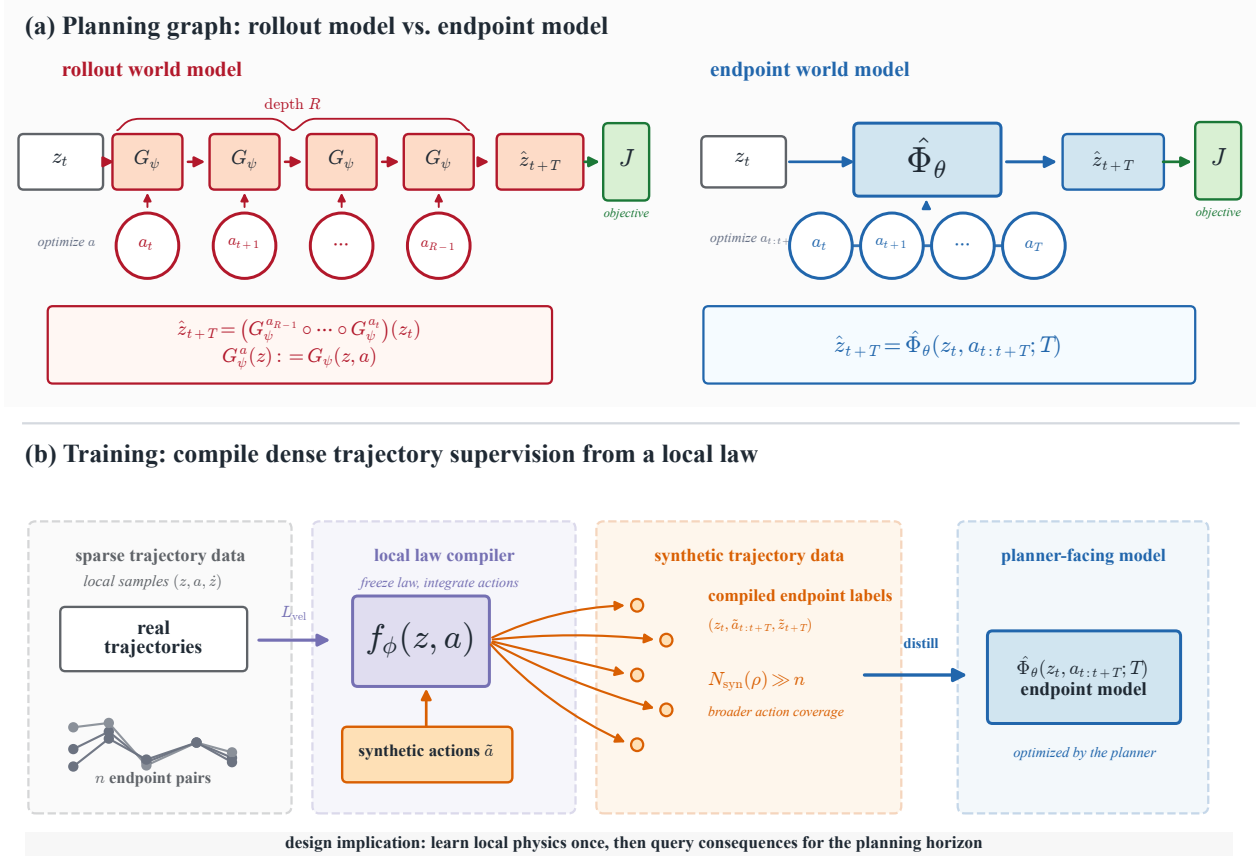


Figure 1. Method overview. (a) **Planning.** The rollout model G_ψ answers a horizon- T query by $R = T/\Delta t$ sequential calls; its action gradient is a product of R Jacobians and is prone to vanishing/exploding (Lemma 3.1). The endpoint world model $\hat{\Phi}_\theta$ maps the initial state and the full action sequence to the endpoint in a single forward pass, so the planner backpropagates through one network Jacobian. The semigroup identity lets the controller query $\hat{\Phi}_\theta$ at a tunable depth, from a single open-loop call to closed-loop replanning. (b) **Two-stage training.** Stage 1 learns a dynamics law f_ϕ from real trajectories by velocity matching. Stage 2 integrates the frozen f_ϕ along synthetic action sequences to compile endpoint targets, which distill the student $\hat{\Phi}_\theta$; the synthetic supervision count $N_{\text{syn}}(\rho) \gg n$ exceeds the paired-data count.

state $z_t := z(t)$ to future state z_{t+T} satisfying

$$\Phi_T^f(z_t, u_{[t, t+T]}) = z_t + \int_t^{t+T} f(z(\tau), u(\tau)) d\tau. \quad (1)$$

A result-driven planner searches for a control signal whose predicted endpoint satisfies the desired condition.

Endpoint World Model. We propose learning a horizon-conditioned proxy for the endpoint operator:

$$\hat{\Phi}_\theta(z_t, u_{[t, t+T]}; T) = \hat{z}_{t+T}, \quad (2)$$

where $\hat{z}_{t+T}$ denotes an estimate of z_{t+T} . Practically, the control signal $u_{[t, t+T]}$ is discretized as an action sequence $a_{0:R-1} \in \mathbb{R}^{Rm}$. Under zero-order hold (ZOH), $u(\tau) = a_j$ for $\tau \in [t + j\Delta t, t + (j+1)\Delta t]$ with $\Delta t = T/R$. Splines and basis expansions give alternative finite-dimensional parameterizations of the same control signal (Hargraves &

Paris, 1987; Kelly, 2017; Ijspeert et al., 2013).

Law-first compilation. We first learn the control law f_ϕ (Chen et al., 2018; Yildiz et al., 2021), then train the endpoint model (2) to match the endpoint operator (1)

$$\Phi_T^{f_\phi}(z_t, u_{[t, t+T]}) := z_t + \int_t^{t+T} f_\phi(z(\tau), u(\tau)) d\tau,$$

where the distillation loss (Salimans & Ho, 2022) is

$$\mathcal{L}_{\text{distill}} = \left\| \hat{\Phi}_\theta(z_t, u_{[t, t+T]}; T) - \Phi_T^{f_\phi}(z_t, u_{[t, t+T]}) \right\|^2. \quad (3)$$

Since f_ϕ maps local state-control pairs to instantaneous state changes, it can be integrated along any well-posed queried control signal $\tilde{u}_{[t, t+T]}$, whether observed, synthetic, or recombined. Each such query defines a teacher target $\tilde{z}_{t+T} = \Phi_T^{f_\phi}(z_t, \tilde{u}_{[t, t+T]})$, yielding dense supervision beyond the finite paired trajectory dataset.

Teacher training. We fit f_ϕ by velocity matching against a five-point central finite-difference estimate of the trajectory velocity. At full data this term suffices. Under data scarcity we also roll f_ϕ forward H steps under constant action and supervise every intermediate state against the observed trajectory; we use $H=100$ (Section B.5). Velocity matching requires constant-action stencil windows. For varying-action datasets, the integrated trajectory loss alone trains f_ϕ (Section B.5).

Student training. We distill $\hat{\Phi}_\theta$ on synthetic action sequences generated at per-step granularity ($R = T/\Delta t$ actions per sample), with action smoothness controlled by a segment count $B \in [1, R]$ sampled *log-uniformly* per sample. Log-uniform B gives each smoothness octave equal training weight, covering both the smooth regime where planning solutions live and the high-frequency regime that the deployment input distribution exercises. Within each segment the action is sampled uniformly from the control range. The teacher then integrates f_ϕ along the resulting R -step action sequence to produce the endpoint target $\hat{z}_{t+T}$.

Optional real-data anchoring. When paired samples $(z_t, u_{[t,t+T]}, T, z_{t+T})$ are available, we regularize $\hat{\Phi}_\theta$ toward the observed true endpoint z_{t+T} rather than the synthetic $\hat{\Phi}_T^f$ induced by f_ϕ (Song & Dhariwal, 2024):

$$\mathcal{L}_{\text{anchor}} = \left\| \hat{\Phi}_\theta(z_t, u_{[t,t+T]}; T) - z_{t+T} \right\|^2. \quad (4)$$

The synthetic targets from f_ϕ supply dense coverage over states, horizons, and control signals, while the real-data anchors accuracy on the support of the paired dataset. Direct endpoint regression is the special case where only this term is used and introduced as a baseline below at (7).

Semigroup-consistency regularization. The endpoint operator induced by any controlled law satisfies the semigroup identity $\Phi_T^f = \Phi_{T-s}^f \circ \Phi_s^f$ for any $0 \leq s \leq T$. When $\hat{\Phi}_\theta$ is queried at intermediate sub-horizons it should compose consistently with its own one-call prediction. We encode this as an auxiliary loss:

$$\mathcal{L}_{\text{sg}} = \left\| \hat{\Phi}_\theta(z_t, u_{[t,t+T]}; T) - \hat{\Phi}_\theta \left(\hat{\Phi}_\theta(z_t, u_{[t,t+s]}; s), u_{[t+s,t+T]}; T-s \right) \right\|^2, \quad (5)$$

with $s \sim \text{Unif}(0, T)$ and T drawn from the same distribution as the distillation loss. This regularizer is optional but useful when the student is trained on a small set of horizon anchors and we want consistent predictions at intermediate horizons not in that set; it also enables the controller to invoke $\hat{\Phi}_\theta$ at variable depth (Section 3).

Average-velocity parameterization. Alternatively, we parameterize the endpoint model as $\hat{\Phi}_\theta(z_t, u_{[t,t+T]}; T) = z_t + T \bar{v}_\theta(z_t, u_{[t,t+T]}; T)$, so the network outputs the average velocity rather than the endpoint itself (Geng

et al., 2026). The distillation loss (3) rescales by $1/T$ to $\bar{v}_\theta(z_t, u_{[t,t+T]}; T) \approx \frac{1}{T} \int_t^{t+T} f_\phi(z(\tau), u(\tau)) d\tau$. This enforces the identity $\hat{\Phi}_\theta(z_t, u_{[t,t+T]}; 0) = z_t$ and keeps the output well-scaled in the small-horizon limit rather than dominated by z_t .

Reference baselines. We justify our proposed world model $\hat{\Phi}_\theta$ against two reference models, one for planning (Section 3) and one for training (Section 4).

The first is a *rollout model* representing local transition (Hafner et al., 2020; Hansen et al., 2022a) $G_\psi(z_t, a_t) \approx z_{t+\Delta t}$. It answers the planner’s finite-horizon query by rolling out for $R = T/\Delta t$ steps:

$$\hat{z}_{t+T} = (G_\psi(\cdot, a_{t+(R-1)\Delta t}) \circ \dots \circ G_\psi(\cdot, a_t))(z_t). \quad (6)$$

The second is an *endpoint regressor*, which shares the endpoint supervision but is directly trained on paired trajectory samples $(z_t, u_{[t,t+T]}, T, z_{t+T})$

$$\mathcal{L}_{\text{end}} = \left\| G_\gamma(z_t, u_{[t,t+T]}; T) - z_{t+T} \right\|^2. \quad (7)$$

3 Planning Advantage

Given a world model, a planner queries predicted future states $\hat{z}_\tau$ to evaluate candidate control signals. Let $r(z, u)$ be an instantaneous reward and $g(z_T)$ a terminal reward. The planner solves the functional optimization

$$\max_{u_{[t,t+T]}} \int_t^{t+T} r(\hat{z}(\tau), u(\tau)) d\tau + g(\hat{z}_{t+T}), \quad (8)$$

to search for the optimal $u_{[t,t+T]}$. Discretize the candidate control signal at the dynamics rate as $a_{0:R-1} \in \mathbb{R}^{Rm}$, where $R = T/\Delta t$. The objective, however, samples only $K \ll R$ reward checkpoints $\tau_j = t + jT/K$, using timestep weight $\Delta\tau = T/K$ and checkpoint action $\bar{a}_j := a_{\lfloor jR/K \rfloor}$. The optimization (8) then becomes finite-dimensional:

$$\max_{a_{0:R-1}} \sum_{j=0}^{K-1} r(\hat{z}_{\tau_j}, \bar{a}_j) \Delta\tau + g(\hat{z}_{t+T}). \quad (9)$$

The planning cost has two factors: (i) how expensively the world model answers each checkpoint query $\hat{z}_{\tau_j}$, and (ii) how large the search space over control signals $u_{[t,t+T]}$ is.

Query Cost. For a single query $\hat{z}_{\tau_j}$, the endpoint model evaluates $\hat{\Phi}_\theta(z_t, u_{[t,\tau_j]}, \tau_j - t)$ once, regardless of the horizon length. The rollout model $G_\psi(z_t, a_t) \approx z_{t+\Delta t}$ returns the same query by $R_j = (\tau_j - t)/\Delta t$ nested calls,

$$\hat{z}_{\tau_j} = (G_\psi(\cdot, a_{R_j-1}) \circ \dots \circ G_\psi(\cdot, a_0))(z_t). \quad (10)$$

For each candidate control signal, the discrete planner in (9) reads K checkpoints $\hat{z}_{\tau_j}$. The rollout model caches intermediate states along the way, so it amortizes to R calls per candidate. The endpoint model directly queries the K required checkpoints. The aggregate per-candidate saving is $R - K$, recovering $R - 1$ in the terminal-only case $K = 1$.

The dynamics rate Δt must be small enough that G_ψ is Markovian and locally learnable at the underlying frame, motor-command, or decision rate. On the dm_control suite this gives $\Delta t \approx 0.02\text{s}$ under the standard action-repeat-2 configuration also used by Dreamer and TD-MPC (Tassa et al., 2018; Hafner et al., 2020; Hansen et al., 2022a), so an episode of $T = 20\text{s}$ already has $R = 10^3$. The number of checkpoints K , however, is set by where reward arrives. Most tasks of interest, e.g., Atari games (Mnih et al., 2015; Bellemare et al., 2013), board games (Schrittwieser et al., 2020), and goal-conditioned robotic manipulation (Andrychowicz et al., 2017), give reward only at score events or task completion, or $K \approx 1$. Hence the saving is essentially the entire rollout length.

Search Structure. For comparison, suppose both models are evaluated on the same grid-level action sequence $a_{0:R-1}$. Let $J_{\hat{\Phi}}(a_{0:R-1})$ denote the planner’s objective (9) under the endpoint model. It depends on this action sequence through one-call queries $\hat{\Phi}_\theta(z_t, u_{[t,\tau_j]}, \tau_j - t)$. Let $J_G(a_{0:R-1})$ denote the same objective under the rollout model. It depends on the action sequence through the depth- R composition (10). Below we show the search structure of $J_{\hat{\Phi}}(a_{0:R-1})$.

Lemma 3.1 (Rollout sensitivities contain sequential Jacobian products). *Let the rollout be $\hat{z}_{t+(r+1)\Delta t} = G_\psi(\hat{z}_{t+r\Delta t}, a_r)$ for $r = [0, R_j - 1]$ and $R_j = (\tau_j - t)/\Delta t$. Assume G_ψ is differentiable. For any $s < R_j$, let $A_\ell := \partial_z G_\psi(\hat{z}_{t+\ell\Delta t}, a_\ell)$ and $B_s := \partial_a G_\psi(\hat{z}_{t+s\Delta t}, a_s)$, then*

$$\frac{\partial \hat{z}_{\tau_j}}{\partial a_s} = A_{R_j-1} A_{R_j-2} \cdots A_{s+1} B_s. \quad (11)$$

Proof. The identity follows by repeated application of the chain rule to the rollout recurrence. $\square$

The product gives horizon-dependent sensitivity bounds. Let $d_{s,j} := R_j - s - 1$ be the number of intervening transition Jacobians $\{A_\ell\}_{\ell=s+1}^{R_j-1}$ between action a_s and checkpoint τ_j . If $\|A_\ell\| \leq L_G$ and $\|B_s\| \leq B_G$ along the rollout,

$$\left\| \frac{\partial \hat{z}_{\tau_j}}{\partial a_s} \right\| \leq B_G L_G^{d_{s,j}}.$$

Thus, when $L_G < 1$, sensitivities to earlier actions decay exponentially with the number of intervening rollout steps. Conversely, if $\sigma_{\min}(A_\ell) \geq \mu_G$ and $\sigma_{\min}(B_s) \geq b_G$, then

$$\sigma_{\min} \left(\frac{\partial \hat{z}_{\tau_j}}{\partial a_s} \right) \geq \left(\prod_{\ell=s+1}^{R_j-1} \sigma_{\min}(A_\ell) \right) \sigma_{\min}(B_s) \geq b_G \mu_G^{d_{s,j}}.$$

So, if all intervening A_ℓ satisfy $\sigma_{\min}(A_\ell) \geq \mu_G > 1$, the sensitivity lower bound grows exponentially with depth.

Corollary 3.2 (Rollout planning requires summing all Jacobian-products). *Suppose the rollout objective J_G follows:*

$$J_G(a_{0:R-1}) = \mathcal{J}(\hat{z}_{\tau_1}(a_{0:R-1}), \dots, \hat{z}_{\tau_K}(a_{0:R-1}), a_{0:R-1}).$$

Then, for any action a_s ,

$$\nabla_{a_s} J_G = \partial_{a_s} \mathcal{J} + \sum_{\substack{j=1 \\ s < R_j}}^K \left(\frac{\partial \hat{z}_{\tau_j}}{\partial a_s} \right)^\top \nabla_{\hat{z}_{\tau_j}} \mathcal{J},$$

where each checkpoint sensitivity $\partial \hat{z}_{\tau_j} / \partial a_s$ requires calculating the sequential Jacobian-products in Lemma 3.1.

Proof. Apply the chain rule to the composition of $\mathcal{J}$ with the rollout-generated checkpoints. $\square$

For any action a_s preceding a queried checkpoint τ_j , i.e., $s < R_j$, Lemma 3.1 shows that $\partial \hat{z}_{\tau_j} / \partial a_s$ contains a product of $R_j - s - 1$ transition Jacobians. Corollary 3.2 further shows that gradients of the rollout objective J_G sum over these checkpoint sensitivities. This is the same failure mode as in recurrent networks (Pascanu et al., 2013; Metz et al., 2021); related mitigations are discussed in Section A.

The endpoint objective bypasses this online recurrent graph by querying each checkpoint directly as $\hat{\Phi}_\theta(z_t, u_{[t,\tau_j]}, \tau_j - t)$. Each checkpoint gradient differentiates through one network Jacobian rather than R_j sequential Jacobians (11). Meanwhile, the K endpoint queries $\{\hat{\Phi}_\theta(z_t, u_{[t,\tau_j]}, \tau_j - t)\}_{j=1}^K$ share the broadcast input $\mathbf{1}_K \otimes z_t$ and can be resolved in one parallel forward pass. Together, these compress the rollout planner’s gradient graph in both depth and width.

Composition depth and replanning are orthogonal.

The semigroup identity $\Phi_T^f = \Phi_{T-s}^f \circ \Phi_s^f$ lets a *single* endpoint query be computed at any composition depth $K_{\text{inf}} \in [1, R]$: partition the queried horizon $T = \sum_{k=1}^{K_{\text{inf}}} \tau_k$ and chain K_{inf} depth-1 calls on the model’s own predicted intermediate state,

$$\begin{aligned} \hat{\Phi}_\theta(z_t, u_{[t,t+T]}; T) &\approx \\ \hat{\Phi}_\theta(\cdot; \tau_{K_{\text{inf}}}) \circ \cdots \circ \hat{\Phi}_\theta(z_t, u; \tau_1), \end{aligned} \quad (12)$$

exact under the identity and tightened by $\mathcal{L}_{\text{sg}}$ (5). This is one open-loop prediction with no environment feedback, sweeping between the depth-1 endpoint query ($K_{\text{inf}}=1$) and the depth- R rollout ($K_{\text{inf}}=R$). Each call is depth-1, so the planner backpropagates through K_{inf} Jacobian terms rather than R (Lemma 3.1 with a product of K_{inf} factors), and composition reaches horizons beyond T_{train} without retraining.

Closed-loop MPC replanning is a separate axis: split the *control* horizon into N_{replan} legs, plan each, execute it on the environment, and replan from the *realized* state. The two are orthogonal and nest multiplicatively – a plan with N_{replan} legs whose per-leg prediction uses depth K_{inf} costs $N_{\text{replan}} K_{\text{inf}} G$ forward calls. Each leg needs accuracy only at its sub-horizon $\leq T_{\text{train}}$; no retraining. Our prediction experiments sweep K_{inf} at $N_{\text{replan}}=1$ (Table 1); our planning

experiments sweep N_{replan} at $K_{\text{inf}}=1$, each leg a single-call sub-horizon prediction (Table 2 and Figure 2). Sub-goal and waypoint planners (Section A) share the depth-1-per-stage property of (12); rollout-within-stage hierarchies (TD-MPC’s H -step lookahead (Hansen et al., 2022b), Dreamer’s imagined rollouts (Hafner et al., 2020)) retain the depth- k Jacobian-product pathology inside each stage. The planning Pareto: $N_{\text{replan}}=1$ wins at $T \leq 2$ s, $N_{\text{replan}}=2$ at $T=3$ s, $N_{\text{replan}}=8$ at $T \geq 4$ s.

4 Training Advantage

Both direct endpoint regression and law-first compilation share the same supervision form. They differ in the target distribution. For $\Delta t_i = T_i/R_i$, let the trajectory dataset be

$$\mathcal{D}_{\text{traj}} = \{(z_0^{(i)}, a_0^{(i)}, z_1^{(i)}, \dots, a_{R_i-1}^{(i)}, z_{R_i}^{(i)}, T_i)\}_{i=1}^n.$$

Direct endpoint regression fits $G_\gamma(z_0^{(i)}, a_{0:R_i-1}^{(i)}; T_i)$ to the observed endpoint $z_{R_i}^{(i)}$ via the tuples $(z_0^{(i)}, a_{0:R_i-1}^{(i)}, T_i, z_{R_i}^{(i)})$. Such supervision must cover the high-dimensional query space of states, horizons, and action sequences.

Law-first compilation generates supervision by integrating a learned control law f_ϕ (1). The law is trained on the $n_{\text{loc}} = \sum_{i=1}^n R_i$ samples $\{(z_r^{(i)}, a_r^{(i)}) \mapsto (z_{r+1}^{(i)} - z_r^{(i)})/\Delta t_i\}_{i,r}$. Given f_ϕ , any well-posed action sequence $\tilde{a}_{0:R_i-1}$ defines a control signal $\tilde{u}_{[0,T_i]}$ and hence a teacher target $\tilde{z}_{R_i} = \Phi_{T_i}^{f_\phi}(z_0^{(i)}, \tilde{u}_{[0,T_i]})$. This target trains the student $\hat{\Phi}_\theta$ via distillation (3). The queried action sequence $\tilde{a}_{0:R_i-1}$ may be observed, synthetic, or recombined, and need not be paired with an observed endpoint. In the following paragraphs, we show law-first compilation dominates direct endpoint regression on $\mathcal{D}_{\text{traj}}$ in two ways: (i) expanding the effective supervision and (ii) shrinking the hypothesis class.

Scale of Synthetic Supervision We now quantify the synthetic supervision for $\hat{\Phi}_\theta$. Let the local training support be $\mathcal{D}_{\text{loc}} = \{(z_r^{(i)}, a_r^{(i)}) : i = 1, \dots, n\}_{r=0}^{R_i-1}$.

Proposition 4.1 (Synthetic supervision count). *Fix $\rho > 0$. For state z , let $\mathcal{A}_\rho(z) := \{a : \text{dist}((z, a), \mathcal{D}_{\text{loc}}) \leq \rho\}$. For anchor $z_0^{(i)}$, let $\tilde{z}_0 := z_0^{(i)}$ and $\tilde{z}_{r+1} := \tilde{z}_r + \Delta t_i f_\phi(\tilde{z}_r, \tilde{a}_r)$. The usable number of synthetic endpoint queries is $N_{\text{syn}}(\rho) := \sum_{i=1}^n |\mathcal{V}_\rho^{(i)}|$, where $\mathcal{V}_\rho^{(i)}$ is a valid action sequence set defined by*

$$\mathcal{V}_\rho^{(i)} := \{\tilde{a}_{0:R_i-1} : \tilde{a}_r \in \mathcal{A}_\rho(\tilde{z}_r) \forall r = 0, \dots, R_i - 1\}.$$

An action sequence is valid as each rollout pair $(\tilde{z}_r, \tilde{a}_r)$ stays within ρ of local training support, so f_ϕ generalizes along the synthetic trajectory and the endpoint target $\tilde{z}_{R_i}$ is reliable. The valid sets $\{\mathcal{V}_\rho^{(i)}\}_{i=1}^n$ induce a mixture over action sequences around the observed anchors. For example, sample an anchor $i \sim \text{Unif}(\{1, \dots, n\})$, then sample $\tilde{a}_{0:R_i-1}$ from a distribution supported on $\mathcal{V}_\rho^{(i)}$. Thus $N_{\text{syn}}(\rho)$ measures the size of the upsampled support for synthetic super-

vision. A ρ -net argument (Vershynin, 2018) on $\mathcal{A}_\rho$ gives $|\mathcal{V}_\rho^{(i)}| \gtrsim (\rho/\epsilon)^{d_a R_i}$ at resolution ϵ , hence $N_{\text{syn}}(\rho)$ grows exponentially in horizon and $N_{\text{syn}}(\rho) \gg n$ whenever $\rho > \epsilon$.

Smaller Hypothesis Class The student endpoint model $\hat{\Phi}_\theta$ distilled from the local teacher f_ϕ shares the same architecture and query space with the endpoint regressor G_γ . The statistical difference originates from the target. Denote $G_\gamma^{a,b}(\cdot) := G_\gamma(\cdot, u_{[a,b]}; b - a)$. Direct endpoint regression learns an unconstrained map $G_\gamma^{t,t+T}(z_t) \approx z_{t+T}$, but the loss does not enforce the semigroup identity

$$G_\gamma^{t,t+T} \neq G_\gamma^{s,t+T} \circ G_\gamma^{t,s}, \quad t < s < t + T.$$

In contrast, law-first compilation generates targets by integrating the learned local f_ϕ . So all teacher labels satisfy

$$\Phi_T^{f_\phi} = \Phi_{t+T-s}^{f_\phi} \circ \Phi_{s-t}^{f_\phi}, \quad t < s < t + T.$$

Thus compilation supplies supervision constrained to a semigroup-consistent target family.

Let $\mathcal{H}_{\text{end}}$ be the hypothesis class of general horizon-conditioned endpoint maps and $\mathcal{H}_{\text{flow}} = \{\Phi^f : f \in \mathcal{F}_{\text{dyn}}\}$ the subclass induced by f . Since integrating f imposes semigroup consistency, $\mathcal{H}_{\text{flow}} \subsetneq \mathcal{H}_{\text{end}}$. Thus law-first compilation restricts the teacher class to $\mathcal{H}_{\text{flow}}$ and trains the student $\hat{\Phi}_\theta$ on semigroup-constrained supervision.

Combining the expanded supervision with ($N_{\text{syn}}(\rho) \gg n$) and shrunk hypothesis class ($\mathcal{H}_{\text{flow}} \subsetneq \mathcal{H}_{\text{end}}$), standard excess-risk bounds (e.g., Rademacher / VC) for law-first compilation are $\mathcal{O}(\text{cap}(\mathcal{H}_{\text{flow}})/N_{\text{syn}}(\rho))$, strictly smaller than $\mathcal{O}(\text{cap}(\mathcal{H}_{\text{end}})/n)$ for direct endpoint regression.

5 Empirical Evidence

We evaluate on Pendulum-v1 (Brockman et al., 2016) with timestep $\Delta t = 0.02$ s. The training set contains 128 trajectories of 300 steps (6 s each); the test set contains 64 trajectories. Section 5.1 reports prediction MSE (Table 1) and terminal-cost planning (Table 2) across $T \in \{1, 2, 3, 4, 5\}$ s on the same model checkpoints and baseline set. The student is distilled on 9 discrete horizon anchors in $[0.96, 4.96]$ s (Section B.7); these evaluation horizons are interpolation queries, only $T=4$ an exact anchor. Section 5.2 and the appendix ablations report prediction MSE over $T = 5.12$ s (256 steps). We compare our endpoint model $\hat{\Phi}_\theta$ trained in (3), against an endpoint regressor G_γ (7) and a rollout model G_ψ fit at Δt and rolled out $R = T/\Delta t$ times (6). Architecture, hyperparameters, planning protocol, and ablations are deferred to Section B.

5.1 Prediction and planning across horizons

The planning task is to drive the pendulum to the upright equilibrium $(\theta, \dot{\theta}) = (0, 0)$ from 30 random initial states. We use $K = 8$ control chunks at every T and report median terminal cost C and the fraction of plans reaching $C < 2$

(success). Endpoint models optimize the action sequence by Adam gradient descent on the planning cost. The rollout model has no usable action gradients (Lemma 3.1), so we plan with cross-entropy method and report an additional gradient-descent ablation row (Section B.3). The MSE protocol is the per-step varying-action input from Table 3, evaluated at each T .

Table 1. Endpoint MSE on per-step varying-action test sequences, $n=256$, deployment-distribution input. Same protocol as Table 3 extended across horizons. Cells: mean / median MSE at the given T . The student is distilled on 9 discrete horizon anchors spanning $[0.96, 4.96]$ s (Section B.7); every T below is an interpolation query (only $T=4$ is an exact anchor). K_{inf} is the open-loop composition depth (the model chained on its own predicted intermediate state, no environment feedback); this is pure prediction, orthogonal to the closed-loop replanning count N_{replan} of Table 2. Bold marks the best mean and best median in each column.

Model	$T=1$	$T=2$	$T=3$	$T=4$	$T=5$
$G_{\gamma}^{\text{scalar}}$	0.85/0.072	2.71/0.19	4.27/0.58	7.97/1.04	8.97/2.08
G_{γ}^{seq}	23.4/17.7	4.42/0.59	25.4/21.1	12.3/4.83	32.5/21.2
G_{ψ} (rollout)	0.51/0.009	3.25/0.044	4.39/0.14	5.43/0.44	6.38/1.06
Ours $\hat{\Phi}_{\theta}$ ($K_{\text{inf}}=1$)	0.36/0.005	1.77/0.027	2.44/0.072	3.22/0.21	5.29/0.82
Ours $\hat{\Phi}_{\theta}$ ($K_{\text{inf}}=2$)	0.34/0.006	2.47/0.015	5.92/1.34	3.53/0.091	5.82/1.01
Ours $\hat{\Phi}_{\theta}$ ($K_{\text{inf}}=4$)	0.59/0.028	2.25/0.035	4.57/0.21	4.39/0.089	5.90/0.24
Ours $\hat{\Phi}_{\theta}$ ($K_{\text{inf}}=8$)	0.52/0.025	2.25/0.21	3.88/0.21	4.47/0.27	5.86/0.75

Table 2. Multi-horizon planning at $T \in \{1, 2, 3, 4, 5\}$ s, $K = 8$. Cells: median terminal cost / reach rate ($C < 2$) over 30 episodes. The student is distilled on 9 discrete horizon anchors spanning $[0.96, 4.96]$ s (Section B.7); every T is an interpolation query (only $T=4$ is an exact anchor). Gradient planners use $G = 50$ Adam steps; sampling planners (CEM and MPPI) use $N_{\text{cand}} = 128$ candidates over $N_{\text{iter}} = 4$ iterations (compute closed form: Table 5). Compute is forward calls per plan; for G_{ψ} it scales as $R = T/\Delta t$ (gradient: $R \cdot G$; CEM/MPPI: $N_{\text{cand}} \cdot N_{\text{iter}} \cdot R$), constant otherwise for one-call endpoint models. One trained $\hat{\Phi}_{\theta}$ underlies the four Ours rows: N_{replan} is the closed-loop MPC replanning count (split the horizon into N_{replan} legs, execute and re-observe between legs; Figure 2), with per-row compute $50 N_{\text{replan}}$. Bold cells mark the best non-oracle median or reach in each column.

Model	Planner	Compute	$T=1$	$T=2$	$T=3$	$T=4$	$T=5$
Oracle	CEM	4,096	1.90/60%	0.89/80%	2.63/37%	4.35/37%	5.60/27%
Oracle	gradient	400	0.96/77%	0.38/97%	0.15/100%	0.12/100%	0.17/100%
$G_{\gamma}^{\text{scalar}}$	gradient	400	1.43/67%	0.94/73%	0.90/70%	0.73/77%	0.67/87%
G_{ψ}	CEM	$\approx 25k \cdot T$	1.65/60%	3.74/47%	3.89/33%	3.59/30%	1.76/57%
G_{ψ}	MPPI	$\approx 25k \cdot T$	2.00/50%	3.23/40%	2.74/40%	1.83/53%	4.34/27%
G_{ψ}	gradient	$\approx 2.5k \cdot T$	0.89/80%	0.52/77%	0.36/70%	0.80/73%	1.26/63%
Ours $\hat{\Phi}_{\theta}$ ($N_{\text{replan}}=1$)	gradient	50	0.86/83%	0.50/93%	0.52/87%	2.63/37%	4.01/27%
Ours $\hat{\Phi}_{\theta}$ ($N_{\text{replan}}=2$)	gradient	100	3.00/30%	1.16/77%	0.084/97%	1.35/57%	0.77/70%
Ours $\hat{\Phi}_{\theta}$ ($N_{\text{replan}}=4$)	gradient	200	5.07/17%	3.15/40%	0.21/80%	1.55/83%	0.27/100%
Ours $\hat{\Phi}_{\theta}$ ($N_{\text{replan}}=8$)	gradient	400	6.65/13%	7.15/13%	2.39/43%	0.009/77%	0.020/100%

Result (MSE, Table 1). Ours wins every column on both mean and median. The best composition depth shifts with T : $K_{\text{inf}}=1$ takes most cells (lowest mean at $T \geq 2$, lowest median at $T \in \{1, 3\}$), $K_{\text{inf}}=2$ takes the $T=1$ mean and $T=2$ median, and $K_{\text{inf}}=4$ takes the long-horizon median at $T \in \{4, 5\}$. G_{γ}^{seq} collapses everywhere: its real-data training distribution is the constant-action trajectory dataset, far from the per-step varying-action test input. Here K_{inf} is *open-loop composition* depth: the model is chained on its own predicted intermediate state, no environment feedback. The composition-depth-vs- T trade-off mirrors the planning table but is not identical. At $T=4, 5$ s composition depth

$K_{\text{inf}}=4$ wins MSE median, but replanning $N_{\text{replan}}=8$ wins planning (Table 2): the planner re-observes the realized state between legs, so it rewards reaching the upright equilibrium, not minimizing open-loop prediction error along arbitrary varying-action sequences. This decoupling motivates reporting both metrics on the same row/column structure.

Result (multi-horizon, Table 2). The same network $\hat{\Phi}_{\theta}$ serves planning at every horizon $T \in \{1, 2, 3, 4, 5\}$ s under closed-loop MPC with replanning count N_{replan} : the horizon is split into N_{replan} legs, each planned over its sub-horizon and executed on the real environment before the next leg replans from the realized state. The N_{replan} that wins grows with T : open-loop ($N_{\text{replan}}=1$) wins both axes against every non-oracle baseline at $T \leq 2$ s, at $48\text{--}96\times$ less compute than G_{ψ} gradient. $N_{\text{replan}}=2$ wins at $T=3$ s ($4\times$ below G_{ψ} on median); $N_{\text{replan}}=8$ wins at $T=4, 5$ s ($89\times$ and $63\times$ below G_{ψ} , at $4\text{--}32\times$ less compute). This recovers the expected replanning-horizon trade-off (Figure 2). The semigroup loss (5) keeps each leg’s sub-horizon prediction consistent, so a single network replans at any depth.

Closed-loop MPC vs G_{ψ} +MPC. The $N_{\text{replan}} > 1$ rows of Table 2 are closed-loop MPC. We compare against G_{ψ} +MPC at the matched $N_{\text{replan}} \in \{1, 2, 4, 8\}$ (Section B.6). At $T=4$ s Ours dominates G_{ψ} +MPC on both axes at $24\times$ less compute; at $T=5$ s G_{ψ} +MPC reaches a $2\times$ lower median only by spending $32\times$ more compute.

Horizon-universal variant. Training the student on continuous-uniform $T \in [0.08, 5.12]$ s in place of the canonical discrete-anchor set yields $\hat{\Phi}_{\theta}^{\dagger}$. It uses the same logB R-step synthetic distribution and reaches median terminal MSE ≤ 0.10 at every $T \in \{0.5, 1, 1.5, 2, 2.5, 3\}$ s on the deployment input distribution. The variant supports arbitrary- T queries and step-level MPC; Section B.7 reports the full grid.

5.2 Data efficiency under scarcity

To test Proposition 4.1 empirically, we retrain each pipeline on 10% of training trajectories (13 of 128). We report held-out endpoint MSE at $T = 5.12$ s on $n = 256$ test sequences with action varying at every env step.

Result (MSE, Table 3). At the canonical $1\times$ budget Ours degrades under scarcity, behind G_{ψ} on the median. Doubling the synthetic-distillation budget halves the median and quadrupling it makes Ours the lowest MSE on *both* mean and median at 10% data. The extra samples are synthetic, drawn from the frozen teacher f_{ϕ} trained on the same 10% real data; no extra real trajectories were observed. $G_{\gamma}^{\text{scalar}}$ cannot represent intra-trajectory action variation through its scalar input and collapses; G_{γ}^{seq} stays above 25 mean at both fractions, its real-data training distribution far from the per-step varying-action test input. Both Ours and G_{ψ} realize $N_{\text{syn}}(\rho) \gg n$ inside the semigroup-consistent class $\mathcal{H}_{\text{flow}} \subsetneq \mathcal{H}_{\text{end}}$, which is why neither collapses on the mean.

Table 3. Endpoint MSE under data scarcity at $T = 5.12$ s on $n = 256$ test sequences with action varying at every env step (deployment-distribution input). Cells: mean / median. For Ours we report the student-budget sweep at 10% data ($1 \times = 48k$, $2 \times = 96k$, $4 \times = 192k$ optimization steps); the canonical 100% recipe runs at $1 \times$ only. The extra steps at $2 \times$ and $4 \times$ draw new *synthetic* samples from the frozen teacher f_ϕ (3); the real-data fraction is unchanged. Teacher recipe and budget rationale: Section B.5.

	100% MSE (mean / med)	10% MSE (mean / med)
$G_{\gamma}^{\text{scalar}}$	9.22/1.65	10.37/6.49
G_{γ}^{seq}	31.84/18.71	25.72/15.80
G_{ψ} (rollout)	6.77/ 0.99	9.11/0.90
Ours $\hat{\Phi}_\theta$ ($1 \times$)	5.63 /1.13	8.57/4.58
Ours $\hat{\Phi}_\theta$ ($2 \times$)	–	6.61/1.84
Ours $\hat{\Phi}_\theta$ ($4 \times$)	–	5.88 / 0.28

Table 4. Planning quality at 10% training data, open-loop ($N_{\text{replan}}=1$). Cells: median terminal cost / reach % over 30 episodes. Forward calls per plan: Ours constant 50; G_{ψ} +gradient scales as $R \cdot G$. Student-budget sweep for Ours uses the same 10% real data with $1 \times / 2 \times / 4 \times$ more *synthetic* distillation samples (no extra real data). Same task and protocol as Table 2.

Method	Calls	$T=2$	$T=3$	$T=5$
$G_{\gamma}^{\text{scalar}}$	400	2.09/43%	3.91/40%	4.73/13%
G_{γ}^{seq}	50	7.77/10%	3.78/27%	5.61/13%
G_{ψ} +gradient	4.8–12.4k	1.38/73%	2.01/50%	2.71/37%
Ours $\hat{\Phi}_\theta$ ($1 \times$)	50	1.72/57%	2.90/30%	4.64/27%
Ours $\hat{\Phi}_\theta$ ($2 \times$)	50	0.96/73%	2.68/47%	4.81/13%
Ours $\hat{\Phi}_\theta$ ($4 \times$)	50	0.72 /80%	1.82 /50%	1.78 /57%

Ours- $4 \times$ recovers full-data mean accuracy (within 5% of Ours- $1 \times$ at 100% data) and improves the median.

Result (planning, Table 4). The budget sweep is monotone at every horizon. $1 \times$ misses the $C < 2$ gate at $T=3$ and $T=5$; $2 \times$ passes $T=2$ only; $4 \times$ passes every cell. At $4 \times$ Ours dominates every non-oracle method on both axes at 96–248 $\times$ less compute than G_{ψ} . The G_{γ} family stays $\geq 2 \times$ worse on median at every cell.

6 Concluding Remarks

We present synthetic data compilation for rollout-free trajectory planning. By replacing recurrent rollout with endpoint consequence queries, the method improves planning cost and action-gradient structure (Section 3). By compiling endpoint targets from a learned dynamics teacher, it improves supervision coverage under paired-data scarcity (Section 4). Experiments support both advantages (Section 5).

References

Andrychowicz, M., Wolski, F., Ray, A., Schneider, J., Fong, R., Welinder, P., McGrew, B., Tobin, J., Pieter Abbeel, O., and Zaremba, W. Hindsight experience replay. *Advances in neural information processing systems*, 30, 2017.

Bellemare, M. G., Naddaf, Y., Veness, J., and Bowling, M.

The arcade learning environment: An evaluation platform for general agents. *Journal of artificial intelligence research*, 47:253–279, 2013.

Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W. Openai gym. *arXiv preprint arXiv:1606.01540*, 2016.

Chane-Sane, E., Schmid, C., and Laptev, I. Goal-conditioned reinforcement learning with imagined subgoals. In *International conference on machine learning*, pp. 1430–1440. PMLR, 2021.

Chen, R. T. Q., Rubanova, Y., Bettencourt, J., and Duvenaud, D. Neural ordinary differential equations. In *Advances in Neural Information Processing Systems*, 2018.

Finn, C. and Levine, S. Deep visual foresight for planning robot motion. In *2017 IEEE international conference on robotics and automation (ICRA)*, pp. 2786–2793. IEEE, 2017.

Geng, Z., Deng, M., Bai, X., Kolter, Z., and He, K. Mean flows for one-step generative modeling. *Advances in Neural Information Processing Systems*, 38:75460–75482, 2026.

Hafner, D., Lillicrap, T., Fischer, I., Villegas, R., Ha, D., Lee, H., and Davidson, J. Learning latent dynamics for planning from pixels. In *International conference on machine learning*, pp. 2555–2565. PMLR, 2019.

Hafner, D., Lillicrap, T., Ba, J., and Norouzi, M. Dream to control: Learning behaviors by latent imagination. In *International Conference on Learning Representations*, 2020.

Hansen, N., Wang, X., and Su, H. Temporal difference learning for model predictive control. In *International Conference on Machine Learning*, 2022a.

Hansen, N., Wang, X., and Su, H. Temporal difference learning for model predictive control. *arXiv preprint arXiv:2203.04955*, 2022b.

Hargraves, C. R. and Paris, S. W. Direct trajectory optimization using nonlinear programming and collocation. *Journal of Guidance, Control, and Dynamics*, 10(4):338–342, 1987. doi: 10.2514/3.20223.

Ho, J., Jain, A., and Abbeel, P. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020.

Hou, X., Huang, X., and Perdikaris, P. Cfo: Learning continuous-time pde dynamics via flow-matched neural operators. *arXiv preprint arXiv:2512.05297*, 2025.

- Ijspeert, A. J., Nakanishi, J., Hoffmann, H., Pastor, P., and Schaal, S. Dynamical movement primitives: Learning attractor models for motor behaviors. *Neural Computation*, 25(2):328–373, 2013. doi: 10.1162/NECO_a.00393.
- Janner, M., Li, Q., and Levine, S. Offline reinforcement learning as one big sequence modeling problem. *Advances in neural information processing systems*, 34: 1273–1286, 2021.
- Janner, M., Du, Y., Tenenbaum, J. B., and Levine, S. Planning with diffusion for flexible behavior synthesis. *arXiv preprint arXiv:2205.09991*, 2022.
- Jurgenson, T., Avner, O., Groshev, E., and Tamar, A. Sub-goal trees a framework for goal-based reinforcement learning. In *International conference on machine learning*, pp. 5020–5030. PMLR, 2020.
- Kelly, M. An introduction to trajectory optimization: How to do your own direct collocation. *SIAM Review*, 59(4): 849–904, 2017. doi: 10.1137/16M1062569.
- Korda, M. and Mezić, I. Linear predictors for nonlinear dynamical systems: Koopman operator meets model predictive control. *Automatica*, 93:149–160, 2018.
- Li, Z., Kovachki, N., Azizzadenesheli, K., Liu, B., Bhattacharya, K., Stuart, A., and Anandkumar, A. Fourier neural operator for parametric partial differential equations. *arXiv preprint arXiv:2010.08895*, 2020.
- Lin, H., Xu, Y.-Y., Sun, Y., Zhang, Z., Li, Y.-C., Jia, C., Ye, J., Zhang, J., and Yu, Y. Any-step dynamics model improves future predictions for online and offline reinforcement learning. *arXiv preprint arXiv:2405.17031*, 2024.
- Lipman, Y., Chen, R. T. Q., Ben-Hamu, H., Nickel, M., and Le, M. Flow matching for generative modeling. In *International Conference on Learning Representations*, 2023.
- Liu, Q. Rectified flow: A marginal preserving approach to optimal transport. *arXiv preprint arXiv:2209.14577*, 2022.
- Lu, L., Jin, P., Pang, G., Zhang, Z., and Karniadakis, G. E. Learning nonlinear operators via deeponet based on the universal approximation theorem of operators. *Nature machine intelligence*, 3(3):218–229, 2021.
- Metz, L., Freeman, C. D., Schoenholz, S. S., and Kachman, T. Gradients are not all you need. *arXiv preprint arXiv:2111.05803*, 2021.
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., et al. Human-level control through deep reinforcement learning. *nature*, 518(7540): 529–533, 2015.
- Pascanu, R., Mikolov, T., and Bengio, Y. On the difficulty of training recurrent neural networks. In *International conference on machine learning*, pp. 1310–1318. Pmlr, 2013.
- Salimans, T. and Ho, J. Progressive distillation for fast sampling of diffusion models. *arXiv preprint arXiv:2202.00512*, 2022.
- Schrittwieser, J., Antonoglou, I., Hubert, T., Simonyan, K., Sifre, L., Schmitt, S., Guez, A., Lockhart, E., Hassabis, D., Graepel, T., et al. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 588(7839): 604–609, 2020.
- Song, Y. and Dhariwal, P. Improved techniques for training consistency models. In *International Conference on Learning Representations*, volume 2024, pp. 15078–15097, 2024.
- Song, Y., Sohl-Dickstein, J., Kingma, D. P., Kumar, A., Ermon, S., and Poole, B. Score-based generative modeling through stochastic differential equations. *arXiv preprint arXiv:2011.13456*, 2020.
- Tassa, Y., Doron, Y., Muldal, A., Erez, T., Li, Y., Casas, D. d. L., Budden, D., Abdolmaleki, A., Merel, J., Lefrancq, A., et al. Deepmind control suite. *arXiv preprint arXiv:1801.00690*, 2018.
- Vershynin, R. *High-dimensional probability: An introduction with applications in data science*, volume 47. Cambridge university press, 2018.
- Yildiz, C., Heinonen, M., and Lahdesmaki, H. Continuous-time model-based reinforcement learning. In *International Conference on Machine Learning*, 2021.

A Related Work and Discussion

Rollout-based MBRL and recurrent planning. Long-horizon rollout depth is a standard difficulty in model-based RL. Backpropagation through a recurrent computation graph produces products of transition Jacobians (Lemma 3.1), which can cause vanishing or exploding sensitivities (Pascanu et al., 2013). Existing world-model methods mitigate the recurrent rollout structure rather than remove it. PlaNet plans by CEM in a learned recurrent latent space, regularized by latent overshooting for multi-step consistency (Hafner et al., 2019). Dreamer learns actors and critics by backpropagating through imagined latent rollouts of a compact recurrent dynamics model (Hafner et al., 2020). TD-MPC truncates planning horizons with a learned terminal value, reducing rollout depth at planning time (Hansen et al., 2022b). All three keep recurrent composition in the model loop. Our formulation removes the recurrent rollout itself. For terminal or sparse-reward objectives, the planner queries $\hat{\Phi}_\theta(z_t, u_{[t, t+\tau]}; \tau)$ directly instead of traversing the recurrent graph (Section 3)

Hierarchical and decomposed planning. Horizon decomposition splits into rollout-within-stage and depth-1-per-stage. Rollout-within-stage schemes shorten the per-stage horizon but each stage is itself a multi-step rollout: TD-MPC’s H -step lookahead (Hansen et al., 2022b), Dreamer’s imagined rollouts (Hafner et al., 2020). Each stage incurs the depth- k Jacobian-product pathology (Lemma 3.1). Depth-1-per-stage schemes replace each stage with a single-shot prediction, as in sub-goal and waypoint reasoners (Jurgenson et al., 2020; Chane-Sane et al., 2021). Our compositional inference (12) is depth-1-per-stage with three differences from those reasoners: each call ingests a full action sub-sequence rather than a goal state (continuous controls, end-to-end gradient); the semigroup loss (5) trains the network for composability; and the composition depth K_{inf} and replanning count N_{replan} are independent controller-time parameters. Diffusion and consistency planners (Janner et al., 2022; Song et al., 2020) are also depth-1-per-stage but compose in generative time, not physical time.

Dynamics and operator learning. A broad line of work learns dynamical laws or solution operators rather than fixed-step transition maps. Neural ODEs and continuous-time MBRL parameterize instantaneous vector fields instead of immediate next-state predictors (Chen et al., 2018; Yildiz et al., 2021). Koopman-based methods learn lifted linear predictors for nonlinear controlled dynamics and have been used inside MPC (Korda & Mezić, 2018). Neural operators, including DeepONet and Fourier neural operators, learn mappings between function spaces and amortize PDE solution operators across initial conditions, parameters, or forcing functions (Lu et al., 2021; Li et al., 2020). These works primarily target prediction, simulation, system identification, or control within a learned or partially specified dynamical system. We use dynamics as a learned representation of controlled state change without a known simulator or symbolic model. The learned local law can then be recombined across supported action sequences to compile finite-horizon endpoint supervision (Sections 2 and 4).

State-change modeling across generative and physical time. Diffusion and score-based models learn denoising scores or reverse-time dynamics and generate samples through iterative solvers over diffusion time (Song et al., 2020; Ho et al., 2020). Flow matching and rectified flow instead learn velocity fields along probability paths (Lipman et al., 2023; Liu, 2022). One-step flow models replace instantaneous velocity with average state change to reduce inference-time solver depth (Geng et al., 2026). Although these methods are usually formulated in generative time, related physical-time work applies flow-matching ideas to continuous-time PDE dynamics by fitting evolution fields rather than autoregressive transitions (Hou et al., 2025). Our work uses the same state-change viewpoint in controlled dynamics, but the output is a planner-facing endpoint predictor rather than a generative sampler (Section 2).

Trajectory prediction and supervision scarcity. Action-conditioned visual foresight, any-step dynamics, and trajectory-level generative models move beyond immediate next-state prediction by modeling longer-horizon futures or full trajectories (Finn & Levine, 2017; Lin et al., 2024; Janner et al., 2021; 2022). Statistically, these approaches learn maps or distributions over high-dimensional trajectory objects whose dimension grows with horizon and action-sequence length. Yet valid trajectories are highly structured: they are consequences of a local controlled law, not arbitrary state-action sequences. Directly fitting future or trajectory maps from paired trajectory data can therefore spend capacity on a large joint object rather than the lower-dimensional dynamical rule that generates it. Law-first compilation uses this structure explicitly by compiling supported synthetic endpoint targets from the learned local law, improving trajectory coverage while constraining targets to dynamics-induced consequences (Section 4 and Proposition 4.1).

B Experimental Setup and Ablations

This section contains architecture, training-recipe, planning-protocol, and ablation details deferred from Section 5. All experiments use the same Pendulum-v1 dataset of 128 training trajectories and 64 held-out test trajectories, each 300 steps

long. Trajectories are generated with a constant uniform control in $[-2, 2]$ and integration step $\Delta t = 0.02$ s. Test trajectories use a different random seed.

B.1 Models

Control law f_ϕ . A multilayer perceptron (MLP) with 4 layers of width 256 and SiLU activations, mapping $(z, u) \mapsto \dot{z}$ with no time input. Trained by velocity matching against 5-point central finite differences (`velocity_matching_central_diff_order: 5`) on the trajectory dataset; the central-difference target reduces the velocity-matching bias from $\mathcal{O}(\Delta t)$ to $\mathcal{O}(\Delta t^4)$. Integrated with RK4 over 8 substeps for distillation queries.

Endpoint regressor G_γ^{scalar} . A multilayer perceptron (MLP) with 4 layers of width 512 and SiLU activations, mapping $(z_t, u, h) \mapsto v$ in the average-velocity parameterization $\hat{z}_{t+h} = z_t + h v$. Trained on observed (z_t, u, h, z_{t+h}) tuples sampled from the training trajectories with h drawn uniformly from $[1, 250]$ environment steps. Assumes the control u is held constant over $[t, t + h]$.

Endpoint regressor G_γ^{seq} . Same attention architecture as $\hat{\Phi}_\theta$ below, but trained on real observed endpoint tuples $(z_t, u_{[t, t+T]}, T, z_{t+T})$ instead of synthetic teacher rollouts. Because the raw trajectories are constant-control, every observed action sequence is degenerate (K copies of one scalar action). Isolates the contribution of synthetic supervision from the contribution of the sequence interface.

Rollout model G_ψ . A multilayer perceptron (MLP) with 3 layers of width 128, mapping $(z_t, u) \mapsto \hat{z}_{t+\Delta t}$ at a fixed step $\Delta t = 0.02$ s. At query time the model is unrolled $R = T/\Delta t = 96$ chunk-aligned base steps under zero-order hold of the candidate control (Section B.3).

Our endpoint world model $\hat{\Phi}_\theta$. A Transformer encoder over the action-chunk tokens $\{(u_r, h_r)\}_{r=0}^{R-1}$ produces a context vector that is concatenated with the initial state z_t and passed to an average-velocity head with 4 layers of width 512. The encoder uses 2 self-attention blocks with hidden dimension 128, 4 heads, and feed-forward width 512. One forward pass produces the predicted endpoint $\hat{z}_{t+T}$ regardless of R .

B.2 Two-stage training

Stage 1 (law). f_ϕ is trained by 5-point central-difference velocity matching on the trajectory dataset. We use batch size 256, 6000 optimization steps, Adam with learning rate 10^{-3} and weight decay 10^{-6} , gradient clipping at 1.0.

Stage 2 ($\hat{\Phi}_\theta$). The student is trained against teacher endpoints $\Phi_T^{f_\phi}(z_t, \tilde{u}_{[t, t+T]})$ obtained by integrating the frozen f_ϕ chunk-by-chunk under synthetic action sequences. We do not use observed endpoints during stage 2. Batch size 256, 6000 optimization steps, Adam with learning rate $7 \cdot 10^{-4}$ and weight decay 10^{-6} , gradient clipping at 1.0.

For each training step, we sample a chunk count K uniformly from $\{2, 4, 8, 16\}$ and a segment count B uniformly from $[1, K]$. The synthetic sequence is piecewise constant in B consecutive segments of equal length. Each segment’s control value is sampled independently from the training-data action pool. The total horizon is fixed to $T = 5.12$ s (256 env steps), so each chunk has duration T/K . The teacher integrates f_ϕ sequentially across the K chunks with 8 RK4 substeps per chunk to produce z_{t+T} .

$B = 1$ recovers the constant-control distribution of the original training trajectories. $B = K$ allows each chunk’s control to differ. Drawing B uniformly per sample exposes the student to a continuum of smoothness levels. The integrated rollout stays near states where f_ϕ is accurate.

B.3 Planning protocol

Gradient planning (endpoint models). For an endpoint operator $\hat{\Phi}_\theta$ and the endpoint regressor G_γ , the action sequence $a = (a_0, \dots, a_{R-1}) \in \mathbb{R}^{Rm}$ is a parameter tensor. Adam optimizes the terminal cost C for 50 steps at learning rate 0.1 with initial standard deviation 0.3. Each step performs one forward and one backward pass through the model. Our endpoint world model $\hat{\Phi}_\theta$ takes the full sequence in one call, so each step costs one forward call (50 total). The endpoint regressor G_γ chains R scalar calls per step ($R \cdot 50$ total).

CEM (rollout baseline). The rollout model G_ψ does not admit a usable action gradient in the recurrent form (Lemma 3.1). We plan with cross-entropy method: a per-chunk Gaussian over the R -chunk control sequence, 128 candidates per iteration, 4 iterations, elite fraction 0.1, initial standard deviation 1.0. Inside each chunk G_ψ unrolls $T/\Delta t = 100$ steps regardless of the planner’s chunk count. Total forward calls per plan is candidates $\times$ iterations $\times$ (R micro-steps per candidate).

MPPI (rollout baseline). We also report a Model-Predictive Path Integral baseline for G_ψ . MPPI uses the same per-chunk Gaussian and the same 128×4 sampling budget as CEM, but updates the Gaussian by the exponentially-weighted mean and variance of the candidate costs with temperature $\lambda = 1.0$ rather than by the top- N elite mean. This is the standard reference planner in recent model-based RL work and rules out the possibility that the rollout-model wedge in Table 2 (the G_ψ CEM/MPPI rows) is an artifact of a weak CEM-style aggregation.

Gradient ablation on G_ψ . For the gradient-descent rows of G_ψ in Table 2, we apply the same Adam recipe through the chained R -step rollout. This pays $R \cdot G$ forward calls per plan and tests whether the gradient pathology of Lemma 3.1 is empirically severe.

Compute counts in Table 2 as a closed form. Forward calls per plan are fully determined by the planner’s parameters and each model’s per-evaluation cost. Let $G = 50$ be the gradient planner’s Adam steps, $N_{\text{cand}} = 128$ and $N_{\text{iter}} = 4$ the CEM candidates per iteration and iteration count, $K = 8$ the number of control chunks, and $R = T/\Delta t$ the env micro-steps over the horizon ($\Delta t = 0.02$ s, so $R = 100$ at $T = 2$ s). One “forward call” is one invocation of `model.predict`, regardless of internal sub-stepping.

A gradient planner makes one forward and one backward pass per Adam step, so per-step cost equals the number of model evaluations required to produce $\hat{z}_T$. Sequence-aware endpoint models ($\hat{\Phi}_\theta$ and G_γ^{seq}) produce $\hat{z}_T$ in one call; scalar endpoint regressors (G_γ^{scalar}) chain K calls; rollout models (G_ψ) chain R calls. CEM evaluates $N_{\text{cand}} \cdot N_{\text{iter}}$ candidates per plan, each requiring as many forward calls as the model needs to score the trajectory. The Oracle is counted at its API granularity (`env.step(chunk_h)` once per chunk).

Table 5. Compute count per plan as a closed form in the planner and model parameters. Substituting $K = 8$, $R = 100$, $G = 50$, $N_{\text{cand}}N_{\text{iter}} = 512$ reproduces the $T = 2$ s column of Table 2; other columns scale with $R = T/\Delta t$.

Model	Planner	Forward calls / plan	Value at $T=2$
$\hat{\Phi}_\theta$ (Ours)	gradient	G	50
G_γ^{scalar}	gradient	$K \cdot G$	400
G_ψ	gradient	$R \cdot G$	5,000
Oracle (true dynamics)	gradient	$K \cdot G$	400
Oracle (true dynamics)	CEM	$N_{\text{cand}}N_{\text{iter}}K$	4,096
G_ψ	CEM	$N_{\text{cand}}N_{\text{iter}}R$	51,200
G_ψ	MPPI	$N_{\text{cand}}N_{\text{iter}}R$	51,200

The closed-form structure makes the compute comparisons in Table 2 reducible to two ratios of model evaluations: (i) *one-call vs. chained-scalar* at the gradient planner ($\hat{\Phi}_\theta$ vs. G_γ^{scalar}): G vs. $K \cdot G$, a factor of $K = 8$; (ii) *gradient vs. CEM* at the same dynamics (Oracle gradient vs. Oracle CEM): KG vs. $N_{\text{cand}}N_{\text{iter}}K$, a factor of $N_{\text{cand}}N_{\text{iter}}/G = 512/50 \approx 10.2$. Composing these, $\hat{\Phi}_\theta$ vs. Oracle CEM is $\approx 81.9\times$ (T-independent, since K is fixed). The first factor is the paper’s contribution (the one-call sequence interface); the second is a known planner-paradigm effect that any gradient-friendly endpoint model would inherit.

Execution and aggregation. After planning, the chosen control sequence is executed on the true environment with zero-order hold inside each chunk. We record the terminal cost $C := \theta^2 + 0.1 \dot{\theta}^2$ at $t = T$. The cost distribution over initial states is bimodal. Plans either reach the target or miss by a wide margin. We aggregate with the median rather than the mean, which would be dominated by the failed plans.

Replanning-depth Pareto sweep. Figure 2 shows the per- T Pareto curves underlying the four Ours rows of Table 2. Each curve sweeps the closed-loop MPC replanning count $N_{\text{replan}} \in \{1, 2, 4, 8\}$ for one fixed planning horizon T . The winning N_{replan} grows with T (open-loop at short T ; depth-8 at long T), bounded by the constraint that the per-leg sub-horizon T/N_{replan} stay inside the student’s training-anchor support.

B.4 Ablations

Smoothness control wins both absolute MSE and data efficiency. The smoothness-controlled student has the lowest in-distribution MSE at both 100% and 20% data, and the lowest scarcity ratio. Independent-uniform synthetic ties absolute MSE at 100% but degrades $1.51\times$ under scarcity.

Sequence input matters at planning time. A scalar-input variant of the law-first student is distilled from f_ϕ but takes one action per call. Its in-distribution MSE sits at 4.92/1.08, close to the independent-uniform sequence variant’s 4.81/1.35;




Figure 2. Closed-loop MPC replanning sweep for the multi-horizon student $\hat{\Phi}_\theta$. Each curve is one planning horizon T , swept over replanning counts $N_{\text{replan}} \in \{1, 2, 4, 8\}$ (compute = $50 N_{\text{replan}}$ forward calls per plan). The Pareto-optimal N_{replan} grows with T : open-loop best at $T \leq 2$ s, depth-2 best at $T = 3$ s (median 0.084), depth-8 at $T \in \{4, 5\}$ s (median 0.009 at $T = 4$, 0.020 at $T = 5$). Replanning hurts at short T where the per-leg sub-horizon T/N_{replan} falls into a poorly-supported region. One trained network underlies every point.

Table 6. Endpoint MSE at $T=2$ s in-distribution on $n = 256$ test sequences with action varying at every env step. Cells: mean / median. Scarcity ratio = 20% mean / 100% mean.

Variant	100% MSE	20% MSE	Scarcity ratio
Scalar distilled (no sequence input)	4.92/1.08	7.77/2.31	1.58
Sequence, independent-uniform synthetic	4.81/1.35	7.26/3.03	1.51
Sequence, smoothness-controlled (ours)	3.86/0.44	3.92/0.73	1.02

gradient planning costs $R \times$ more forward calls per Adam step than the one-call sequence student, losing the compute advantage of Section 5.1.

B.5 Scarcity-recipe ablation

Velocity matching alone preserves a sharp typical-trajectory error (low median) but admits a heavy off-manifold tail at sparse data. Endpoint loss at large supervision horizon diverges from coarse-substep integration; at short horizon it destroys the median. The integrated trajectory loss is stable across horizons and monotone in H on the mean: longer supervision cuts the tail. $H=150$ hits a stability ceiling.

The integrated loss assumes zero-order hold (ZOH) over the H -step rollout window: the action is sampled once at t_0 and reused for all H chained one-step predictions. The training data is generated with action-hold windows of 300 steps ($\gg H$), so every sample window we draw is constant-action by construction. The supervision target z_{t_0+h} matches the integration under the same constant action that drove the data; the loss has no action-mismatch bias.

TFH (Trajectory-Force, dense parallel per-step supervision evaluated at observed (z, u) tuples) tests whether on-manifold supervision can replace the chained $H=100$ loss. TFH alone has the lowest median (0.081) but the mean stays at 7.41, above the recipe’s 4.86. vm+TFH sharpens median (0.096) but the mean climbs to 13.01. Adding chained- $H=20$ to TFH makes things worse on both axes. Only long-chained supervision ($H=100$) reaches a low mean. TFH and chained supervision target different failure modes: TFH supervises at observed (z, u) tuples and sharpens the typical case; chained supervision visits off-manifold states the model drifts into during training and catches the mean tail. The scarcity heavy-tail problem is off-manifold; only chained supervision resolves it.

For varying-action datasets, the cen-5 velocity-matching term can be replaced by the integrated trajectory loss alone. At full

Table 7. Teacher step-by-step rollout MSE at $T = 5.12$ s on $n = 256$ test sequences with action varying at every env step (Table 3 probe), 10% training data. Mean and median are reported separately to expose the heavy-tail asymmetry.

Teacher loss	Mean	Median
G_ψ -10% (rollout baseline)	9.11	0.90
vm-only (cen-5)	17.97	0.210
cen-3 vm-only	12.58	0.222
vm + endpoint loss ($H=50$)	7.61	0.433
vm + integrated ($H=10$)	12.89	0.112
vm + integrated ($H=50$)	8.84	0.579
vm + integrated ($H=100$, recipe)	4.86	0.117
vm + integrated ($H=150$)	6.85	0.445
TFH=100 (no vm)	7.41	0.081
TFH=100 + chained- $H=20$ (no vm)	14.79	6.82
vm + TFH=100	13.01	0.096
vm + TFH=100 + chained- $H=20$	17.48	0.218
vm + TFH=100 + chained- $H=100$	6.44	0.785
integrated alone ($H=1$, no vm)	23.30	0.140
integrated alone ($H=100$, no vm)	7.74	1.56

data this swap is neutral or better: integrated-alone $H=100$ reaches teacher rollout mean 3.76 on the constant-action probe, matching vm-only’s 2.36 and vm+integrated- $H=100$ ’s 4.70. Under scarcity the cen-5 anchor adds a $1.6\times$ polish on the mean (integrated-alone $H=100$ at 10%: 7.74 vs vm+integrated- $H=100$ ’s 4.86), but the chained integrated loss carries the structural lift relative to vm-only’s 17.97. The recipe choice is therefore between cen-5 + integrated (when the training data has constant-action segments long enough for the stencil) and integrated alone (for arbitrary-action data).

Student step budget under scarcity. Per-batch teacher signal is independent of real-data count: logB synthetic queries draw fresh $(z, u_{[t, t+T]})$ pairs and integrate the frozen f_ϕ to produce endpoints. The student needs more optimization steps under scarcity to absorb this fixed distribution. The budget sweep is in Tables 3 and 4; $4\times$ (192k steps) is the canonical scarcity recipe. The compute cost is one-time training, paid against the constant 50-call inference budget at deployment.

B.6 Closed-loop MPC evaluation

Open-loop planning ($N_{\text{replan}}=1$) fixes the action sequence at $t=0$ and executes the full T -step plan. MPC splits T into N_{replan} legs, executing each on the real environment and replanning the remaining horizon from the realized state. This is replanning, not composition: each leg is an independent planning problem over its sub-horizon T/N_{replan} , so the model predicts only that section. We sweep $N_{\text{replan}} \in \{1, 2, 4, 8\}$ for Ours and G_ψ +gradient. Same task as Table 2: 30 episodes, gradient planner with 50 Adam steps per replan, PYTHONHASHSEED=0.

Table 8. Closed-loop MPC sweep across replanning counts and planning horizons. Cells: median terminal cost / reach % / total forward calls per plan. Both methods replan on the identical N_{replan} legs; each leg plans its sub-horizon with 50 Adam steps. G_ψ +MPC rolls out the remaining horizon at each leg and pays $\approx R \cdot G$ calls per leg where R is the remaining-horizon length, versus 50 for Ours.

Method, T	$N_{\text{replan}}=1$	$N_{\text{replan}}=2$	$N_{\text{replan}}=4$	$N_{\text{replan}}=8$
Ours, $T=2$	0.50/93/50	1.16/77/100	3.15/40/200	7.15/13/400
Ours, $T=3$	0.52/87/50	0.084/97/100	0.21/80/200	2.39/43/400
Ours, $T=4$	2.63/37/50	1.35/57/100	1.55/83/200	0.009/77/400
Ours, $T=5$	4.01/27/50	0.77/70/100	0.27/100/200	0.020/100/400
G_ψ , $T=2$	0.52/77/4800	0.90/77/4800	2.86/40/4800	3.34/20/6400
G_ψ , $T=3$	0.46/70/7600	0.24/97/7200	0.03/93/8000	3.18/23/6400
G_ψ , $T=4$	0.49/73/10000	2.75/47/9600	2.59/40/9600	0.41/73/9600
G_ψ , $T=5$	0.87/67/12400	1.49/53/12800	0.19/97/12800	0.009/100/12800

For Ours the optimal N_{replan} grows with T : open-loop at $T=2$, depth-2 at $T=3$, depth-8 at $T=4$ and $T=5$. G_ψ +MPC’s optimum is less monotone and its per-replan cost is $96\text{--}256\times$ ours. Main-text headline cells: Section 5.1.

B.7 Horizon-universal student

The canonical Ours student is distilled at 9 discrete horizon anchors $T \in \{0.96, 1.44, 1.92, 2.40, 3.04, 3.52, 4.00, 4.48, 4.96\}$ s. The horizon-universal $\hat{\Phi}_\theta^\dagger$ variant replaces the anchor distribution with continuous-uniform $T \in [0.08, 5.12]$ s; the synthetic action distribution (logB R-step), architecture, optimizer, budget, and teacher are identical.

Table 9. Endpoint MSE (mean / median) for $\hat{\Phi}_\theta^\dagger$ across T on $n = 256$ per-step varying-action test sequences.

T (s)	0.5	1	1.5	2	2.5	3	5
$\hat{\Phi}_\theta^\dagger$ mean	0.31	0.26	1.42	1.89	2.71	2.59	5.07
$\hat{\Phi}_\theta^\dagger$ med	0.003	0.005	0.10	0.019	0.096	0.046	1.11

$\hat{\Phi}_\theta^\dagger$ is valid at any $T \in [0.08, 5.12]$ s. At $T = 0.5$ s it reaches median 0.003, $25\times$ lower than the canonical Ours which extrapolates below its shortest anchor (0.075 median). Aggregate per-step MSE (mean over $T \in \{0.5, 1, \dots, 5\}$ s): $\hat{\Phi}_\theta^\dagger$ 2.23, canonical Ours 1.79.

B.8 Student recipe ablation

We ablate the student synthetic-action distribution on per-step varying-action MSE at $T = 5.12$ s ($n = 256$). Each row is a separate training run; same teacher (vm-only at 100% data) and 48k-step budget. The K-chunked-anchor variant samples K discrete control chunks per sequence (constant action per chunk); R-step variants sample R per-step actions per sequence with segment count B controlling smoothness.

Table 10. Student recipe ablation: per-step varying-action MSE at $T = 5.12$ s.

Student recipe	MSE mean	MSE median
K-chunked anchors ($K \in \{8, 16, 32, 64\}$)	324.4	151.5
R-step bucket ($R \in \{8, \dots, 256\}, B=R$)	7.69	3.03
R-step freeA-extB (uniform $B \in [1, R]$)	4.76	0.63
R-step logB (canonical)	5.63	1.13

K-chunked anchor training fails by $60\times$ on per-step varying-action input. R-step variants fix it. Uniform- B freeA-extB has the lowest MSE but loses planning at open-loop short- T ($T=2$ $N_{\text{replan}}=1$ median 4.27 vs canonical 0.50): uniform- B overweights high-frequency segments, leaving the smooth-piecewise regime where planning solutions live underrepresented. Log-uniform B gives each smoothness octave equal weight ($\sim 37\%$ of samples at $B \leq 16$ vs $\sim 6\%$ under uniform- B), accurate at both extremes.

B.9 Generalization to unseen horizons

We test whether the two-stage construction transfers beyond the training horizon range. The first two rows hard-cap training data at $h \leq 50$ steps; the third keeps the same f_ϕ but trains the student on synthetic teacher rollouts up to $h \leq 256$, well beyond the raw-data limit. Test sequences have action varying at every env step.

Table 11. Endpoint MSE (mean) at in- and out-of-distribution horizons on $n = 256$ test sequences with action varying at every env step. Training horizons are capped at $h = 50$. Entries marked * are unseen at training time.

Endpoint MSE	$h=50$	$h=75^*$	$h=100^*$
$G_\gamma, h \leq 50$ (raw-data limit)	0.59	11.82	18.27
Ours $\hat{\Phi}_\theta, h \leq 50$ synthetic	2.40	23.43	62.98
Ours $\hat{\Phi}_\theta, h \leq 256$ synthetic	0.45	1.35	2.04

The $h \leq 50$ students fail to generalize past their training cap: G_γ 's MSE rises from 0.59 at $h=50$ to 18.27 at $h=100$; the $h \leq 50$ Ours student degrades to 62.98. The $h \leq 256$ Ours student keeps MSE under 2.1 across all three horizons, $9\times$ lower than G_γ and $31\times$ lower than the $h \leq 50$ Ours student at $h=100$. Horizon extrapolation requires using f_ϕ to manufacture supervision past the raw-data range.